FLORIDA POLYTECHNIC UNIVERSITY

Neo4J

Intro

What is Neo4J

- Worlds leading graph database
- Fast traversal of nodes and relationships
- Four million hops per second
- Can store trillions of entities
- Cypher graph query language
- Scalable fault tolerant clusters

History

- Founded in 2007
- 1.0 Released in 2010
  - 34 billion node and relationship limit
- 2.0 released in 2013
- 3.0 released in 2016
  - 36m in funding secured
  - Unlimited nodes

FLORIDA POLYTECHNIC UNIVERSITY

Where does it fit in CAP Theorem

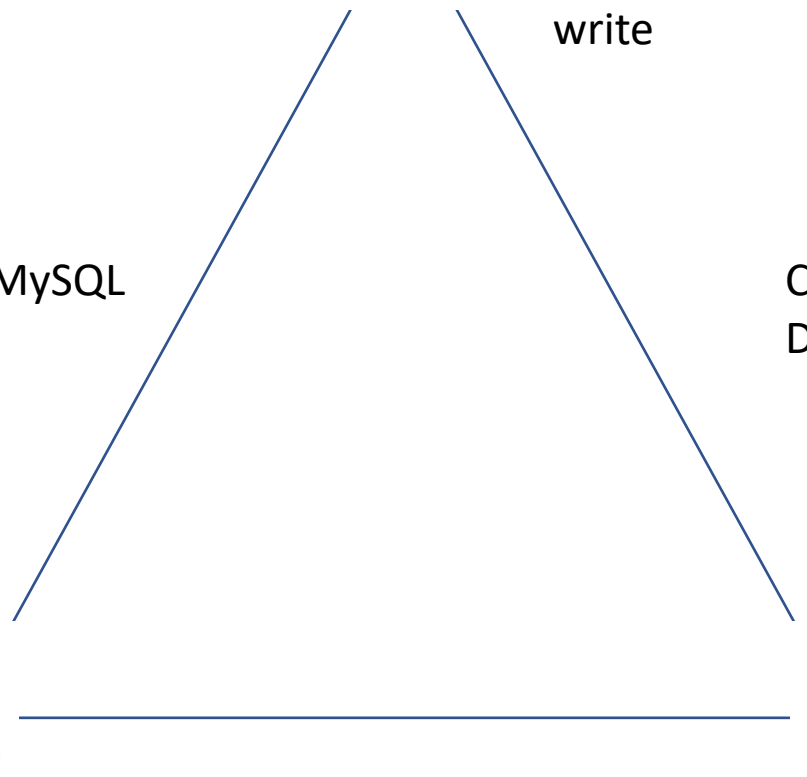

Availability:
Each client can always read and write

RDBMS, Postgres, MySQL
Neo4J

CouchDB, Cassandra, DynamoDB**,** Riak

Consistency: All clients always have the same view of data

Big Table, Hbase, MongoDB, Redis

Partition Tolerance: Works well despite physical network partitions

FLORIDA POLYTECHNIC UNIVERSITY

How is Neo4J Different

- Graph Storage
- Graph Processing Engine
- Relationships matter more than data points

Neo4J Architecture Core Servers

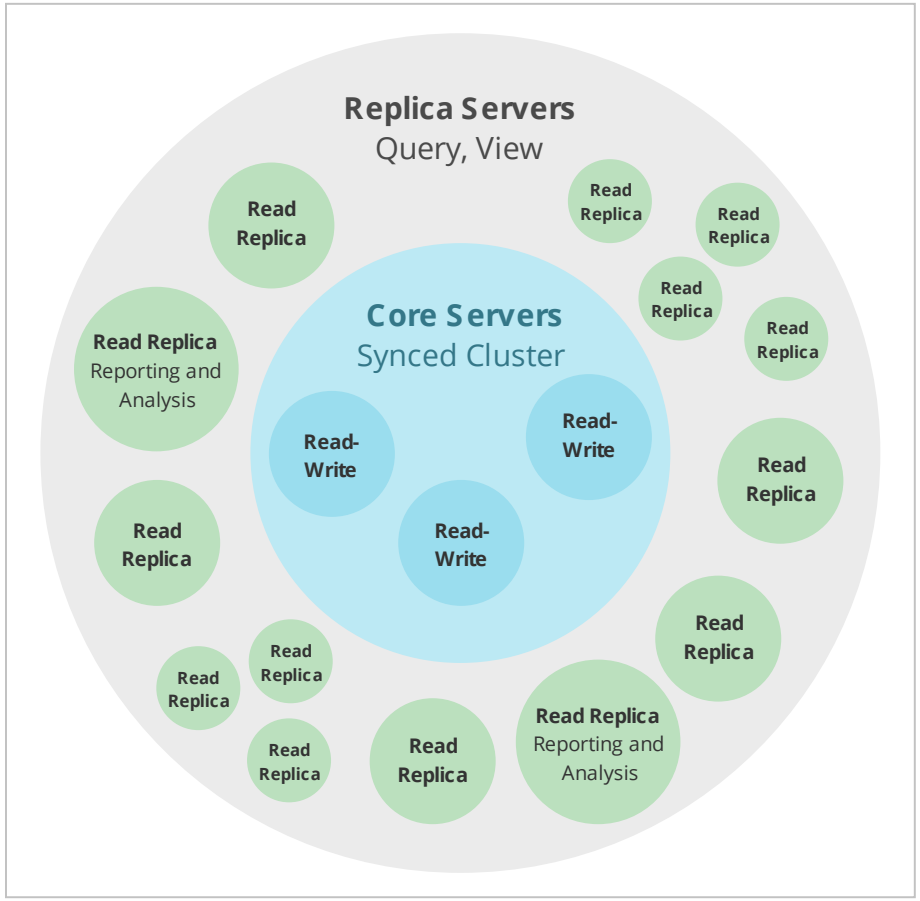

FLORIDA POLYTECHNIC UNIVERSITY

Neo4J Architecture Read Replicas

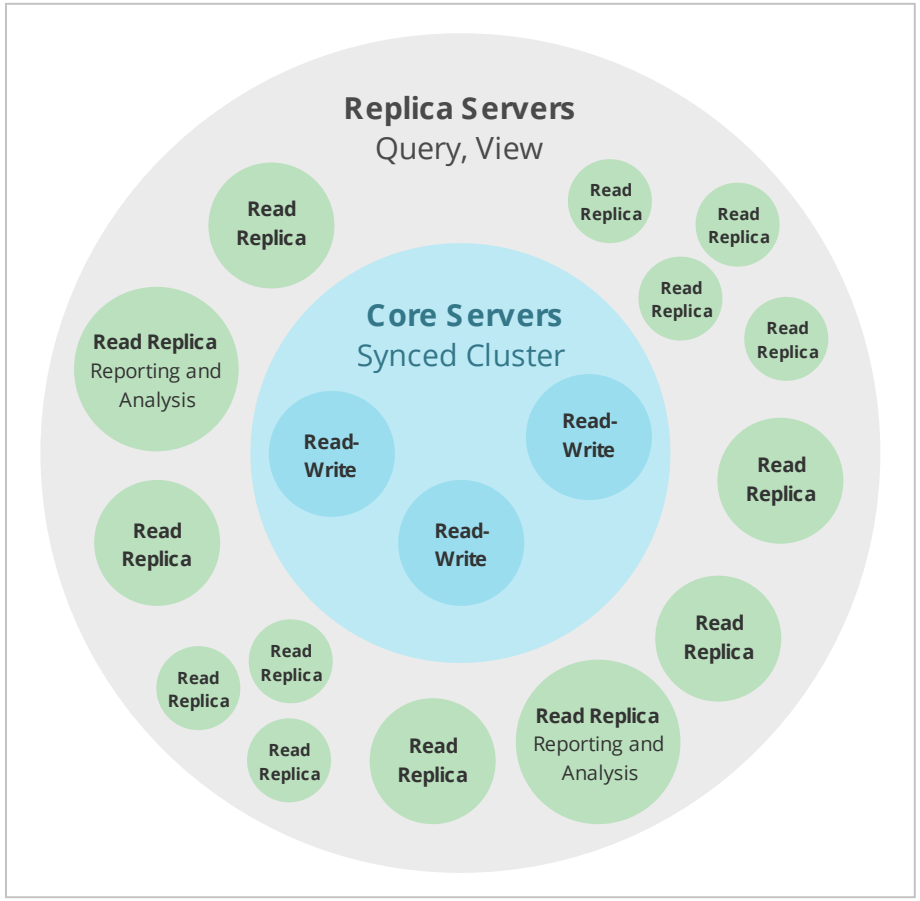

FLORIDA POLYTECHNIC UNIVERSITY

System Architecture

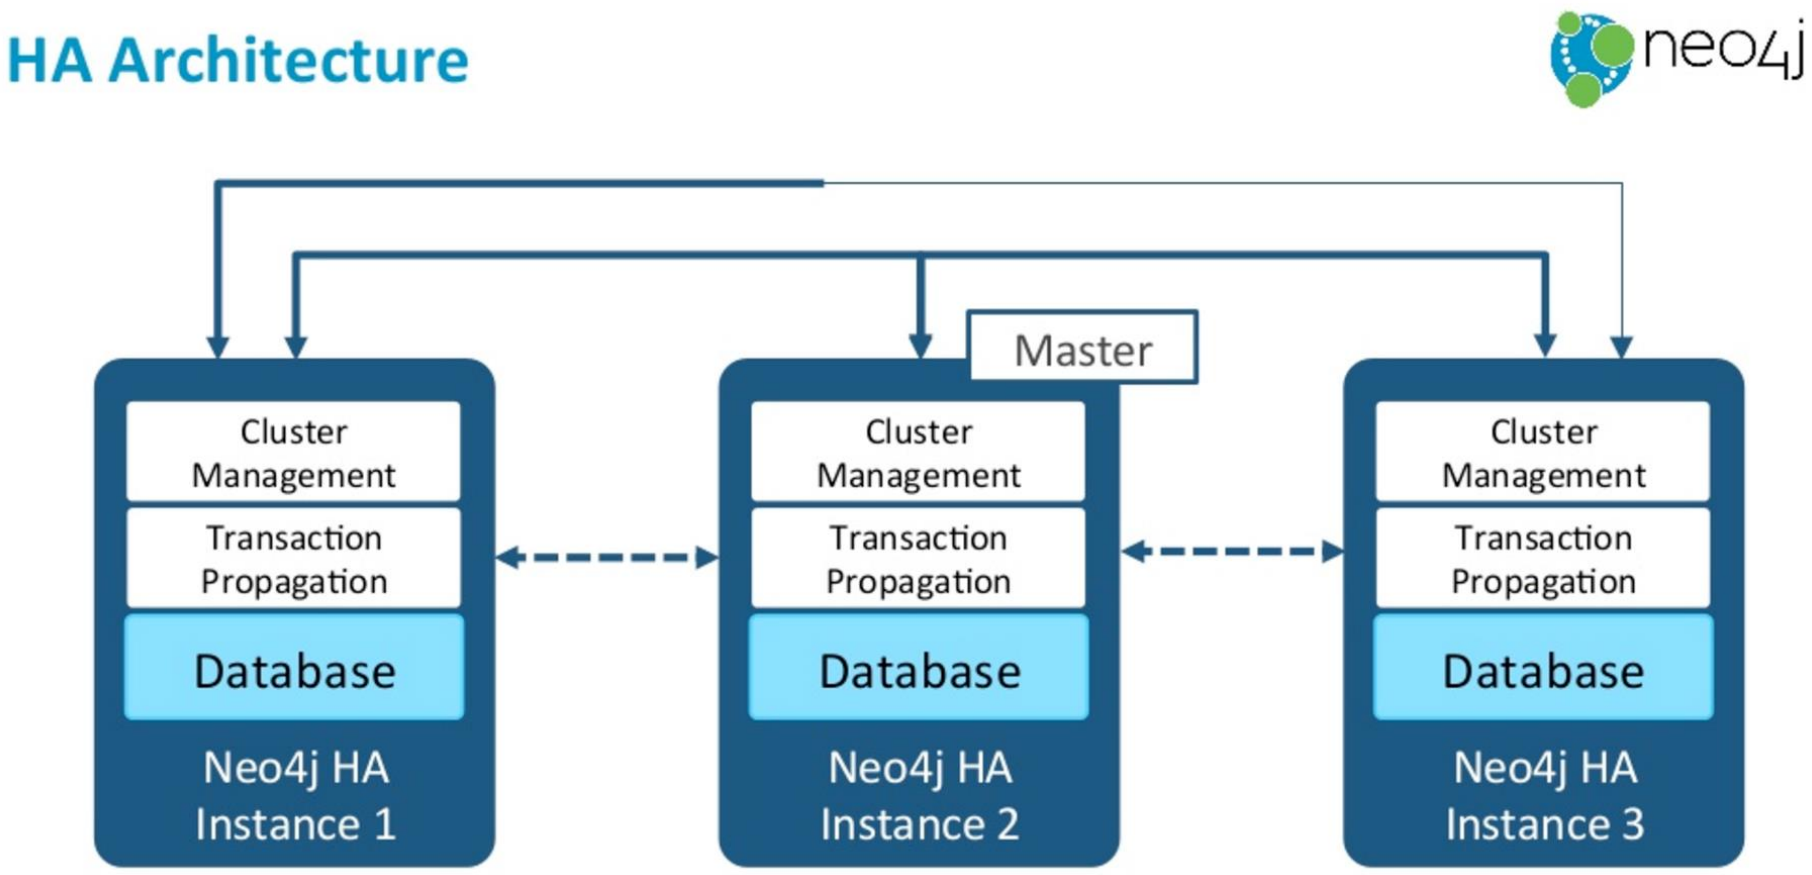

Load Balancing for Read Throughput

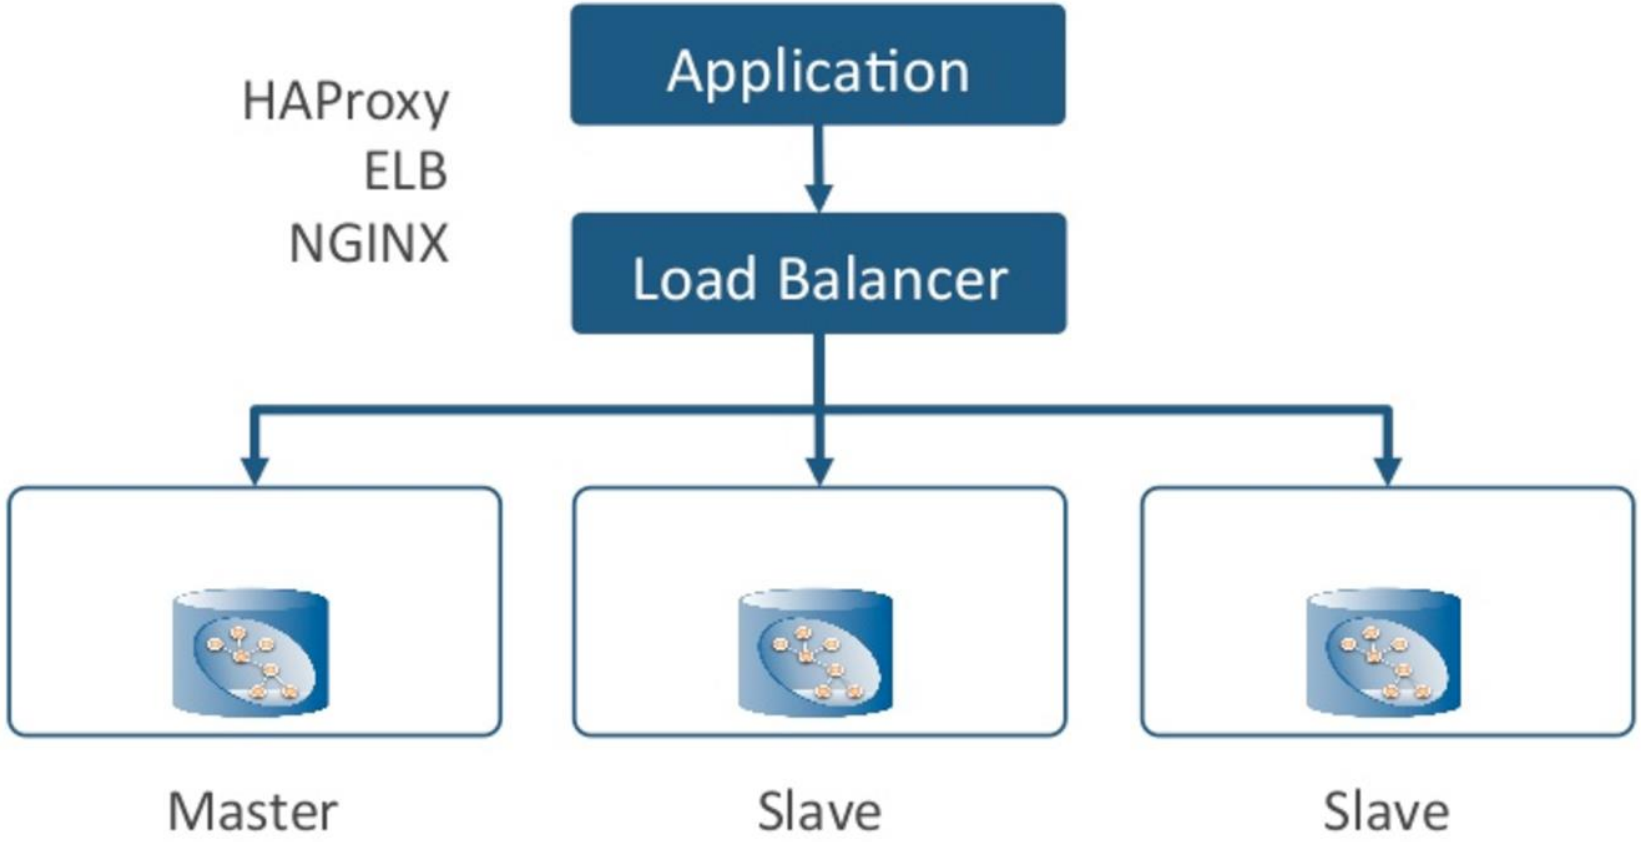

Load Balance for Writes

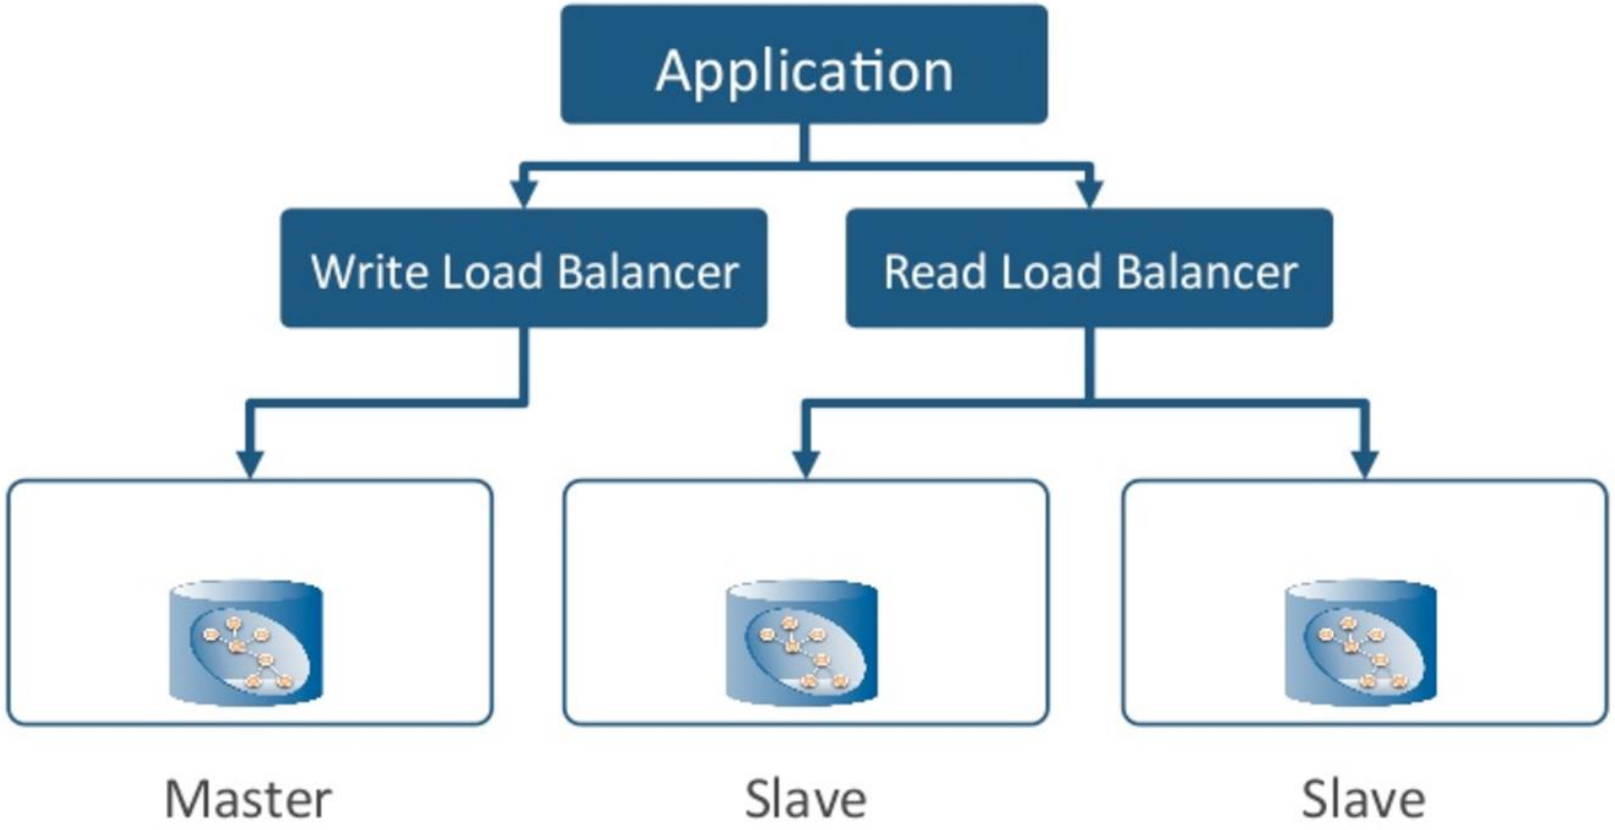

Causal Consistenty

Clustering – Discovery Protocol

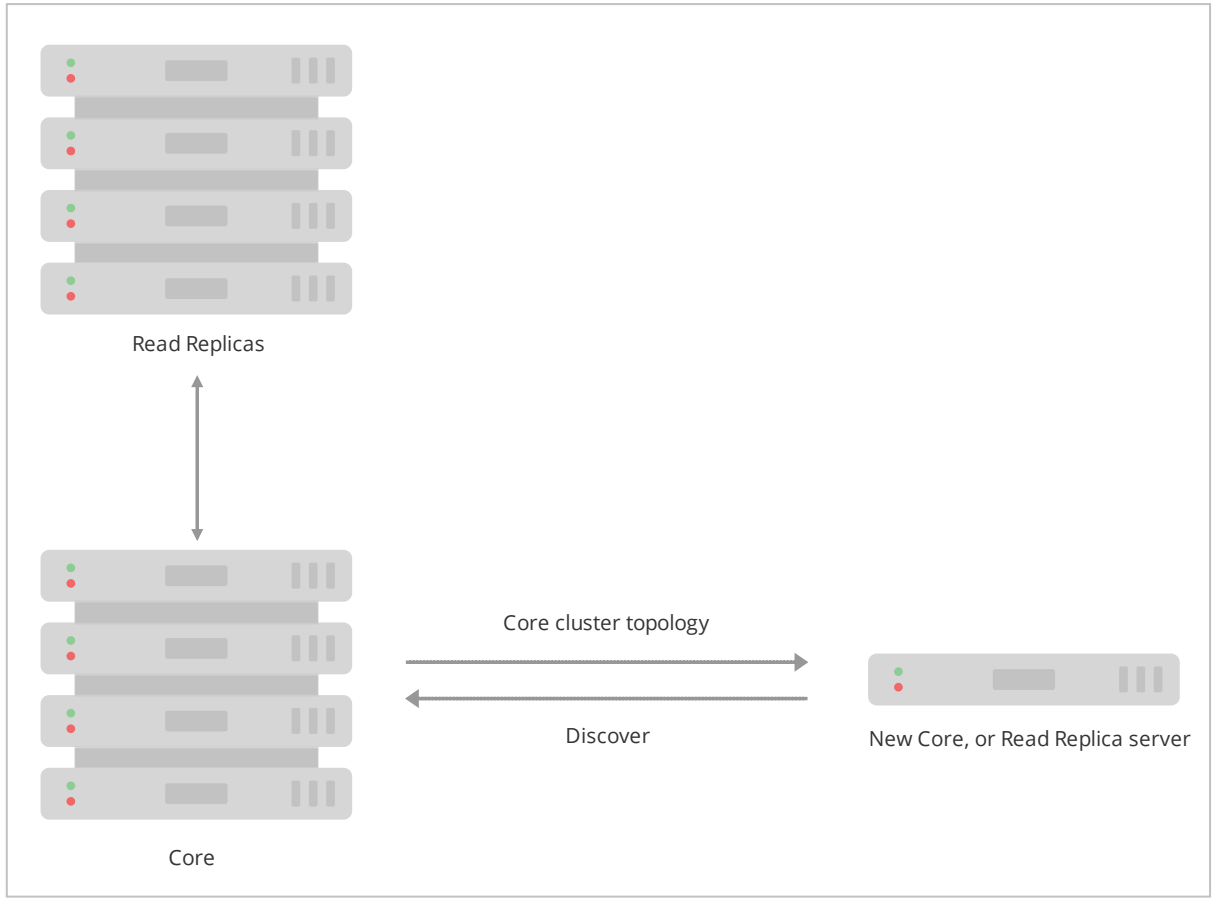

Clustering – Read Replica membership

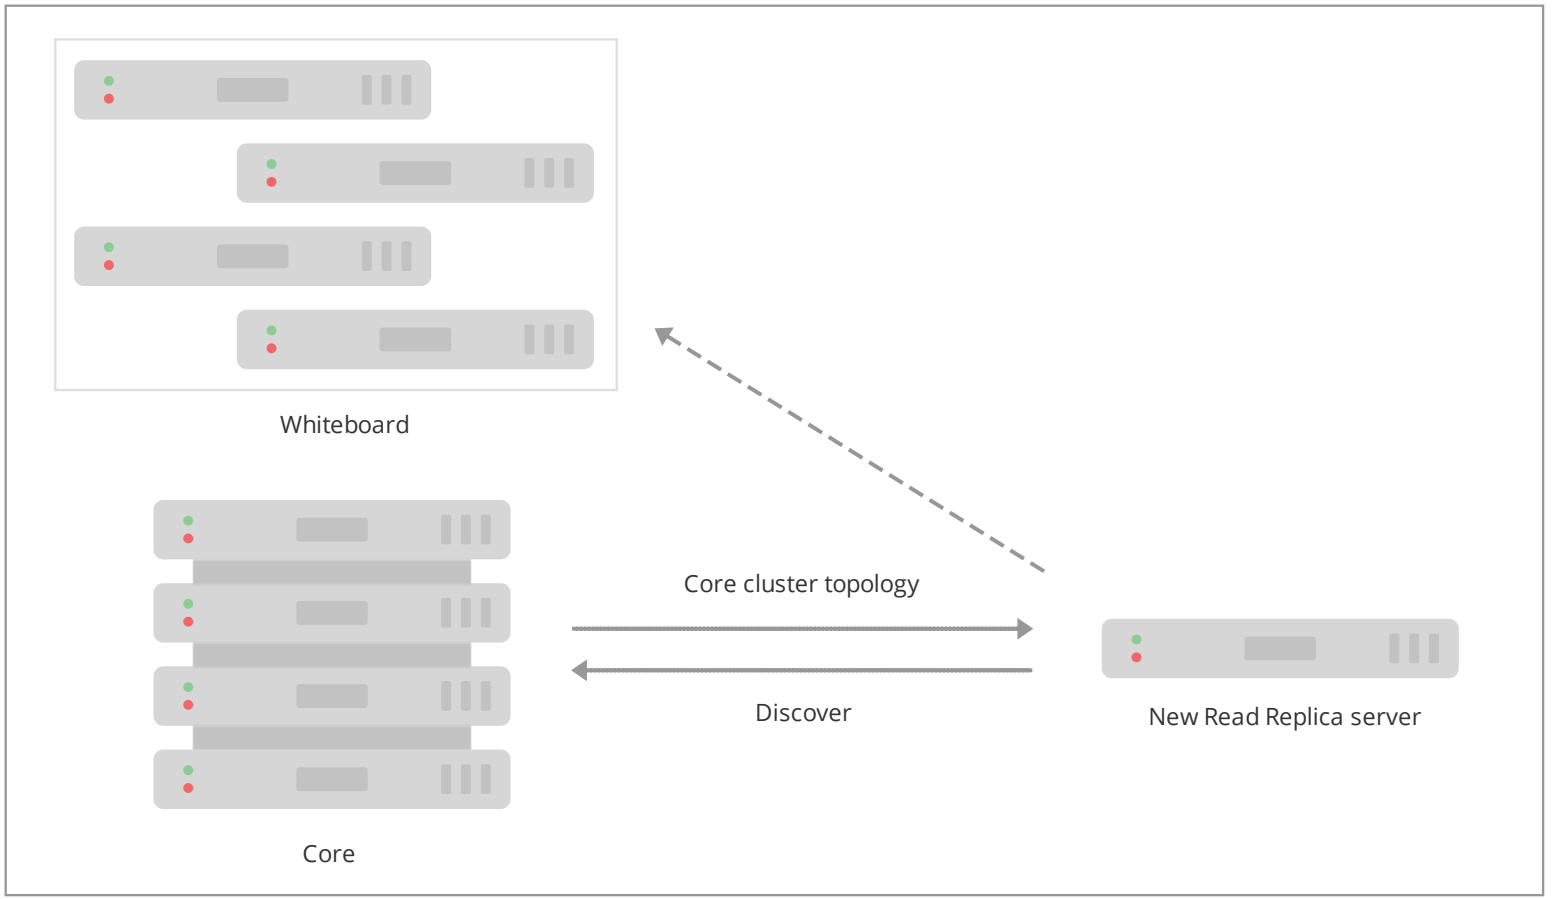

Clustering Read Replica

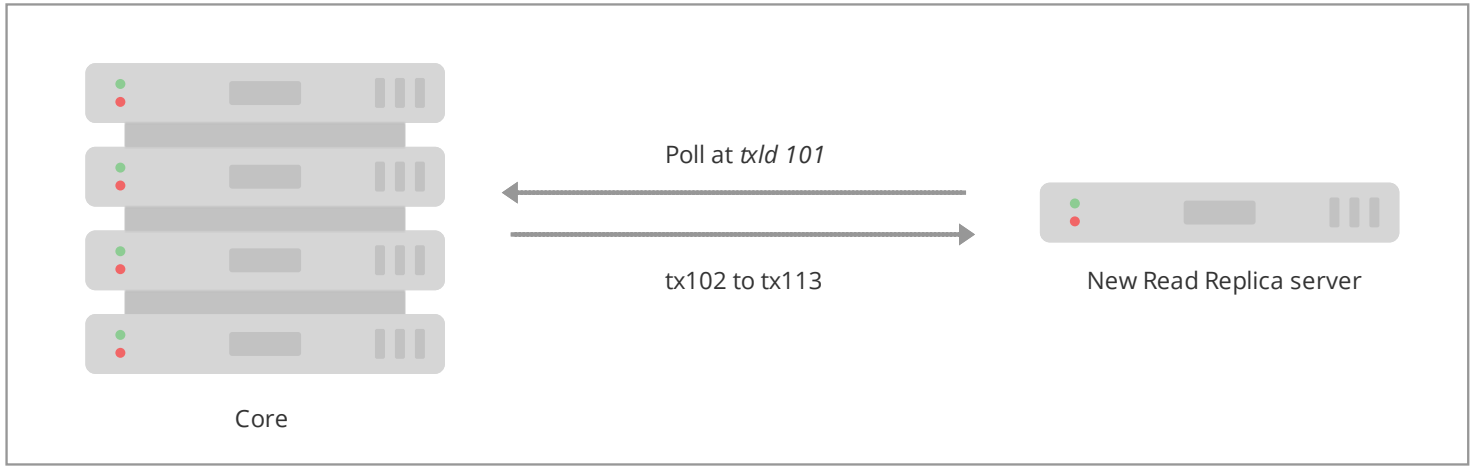

Server Architecture

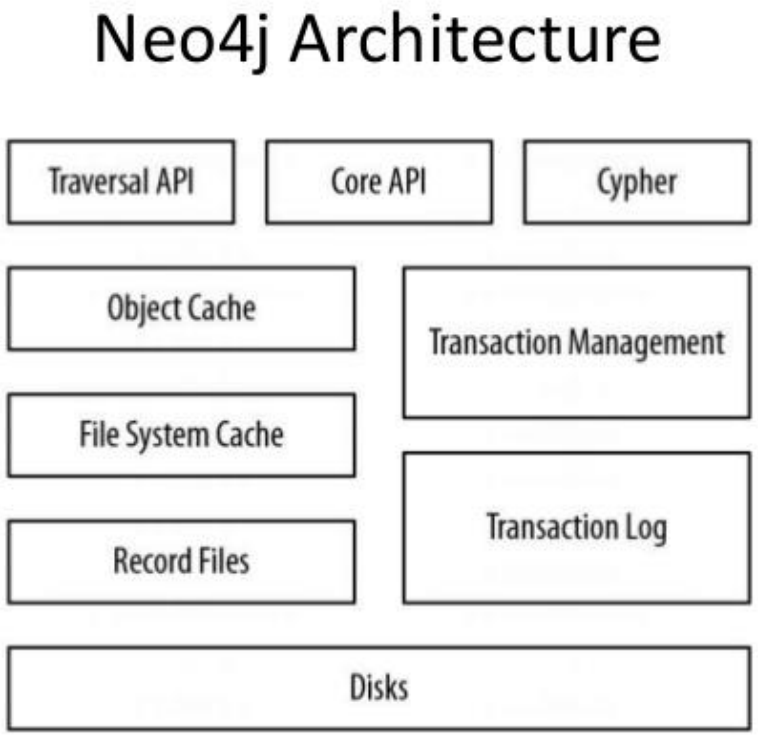

CQL

- Cypher Query Language is Neo4j's open graph query language

ACID Compliant

- Database operations that access the graph, indexes, or the schema must be performed in a transaction.

Memory Configuration

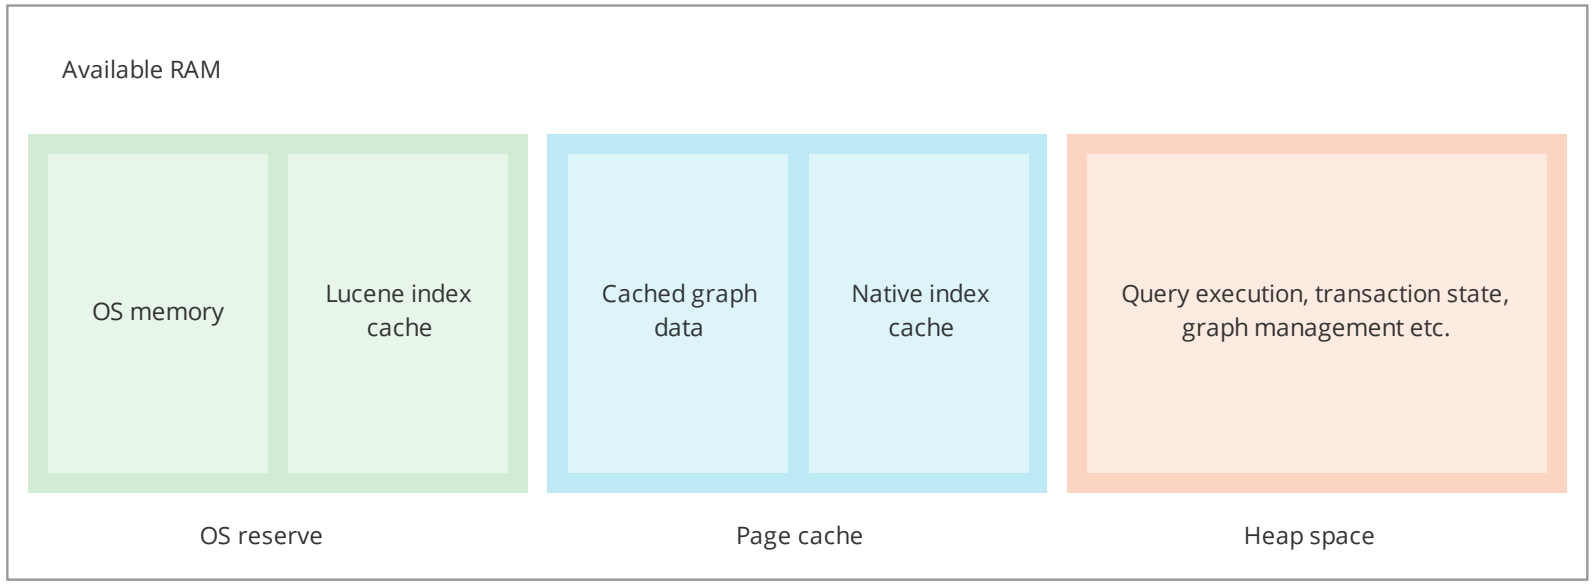

Traversal

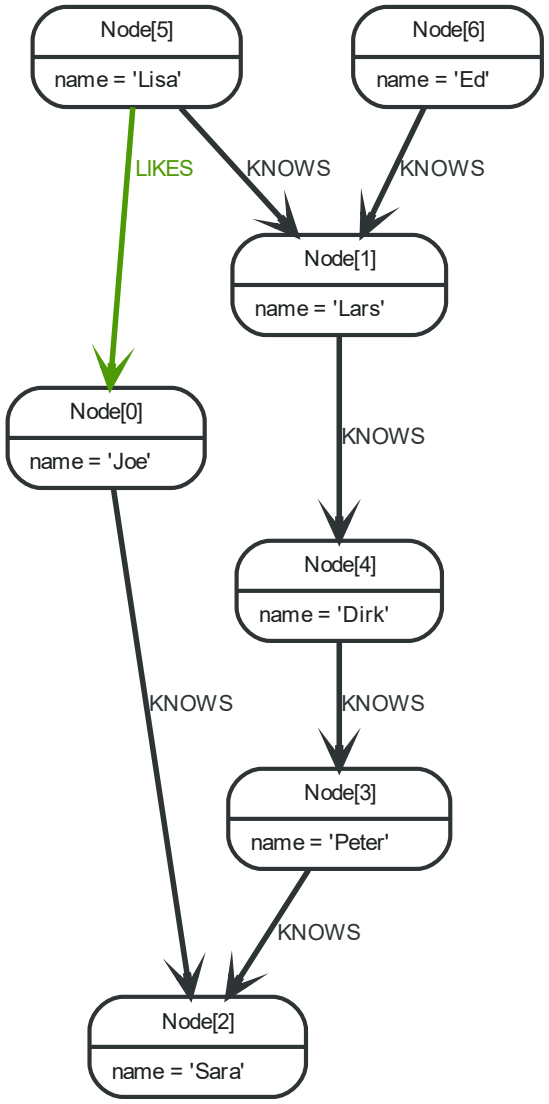

With the definition of the `RelationshipTypes` as

```
private enum Rels implements RelationshipType
{
    LIKES, KNOWS
}
```

The graph can be traversed with for example the following traverser, starting at the ``Joe'' node:

```
for ( Path position : db.traversalDescription()
        .depthFirst()
        .relationships( Rels.KNOWS )
        .relationships( Rels.LIKES, Direction.INCOMING )
        .evaluator( Evaluators.toDepth( 5 ) )
        .traverse( node ) )
{
    output += position + "\n";
}
```

The traversal will output:

```
(0)
(0)<-[LIKES,1]-(5)
(0)<-[LIKES,1]-(5)-[KNOWS,6]->(1)
(0)<-[LIKES,1]-(5)-[KNOWS,6]->(1)<-[KNOWS,5]-(6)
(0)<-[LIKES,1]-(5)-[KNOWS,6]->(1)-[KNOWS,4]->(4)
(0)<-[LIKES,1]-(5)-[KNOWS,6]->(1)-[KNOWS,4]->(4)-[KNOWS,3]->(3)
(0)<-[LIKES,1]-(5)-[KNOWS,6]->(1)-[KNOWS,4]->(4)-[KNOWS,3]->(3)-[KNOWS,2]->(2)
```

FLORIDA POLYTECHNIC UNIVERSITY

Data Modeling – Whiteboarding

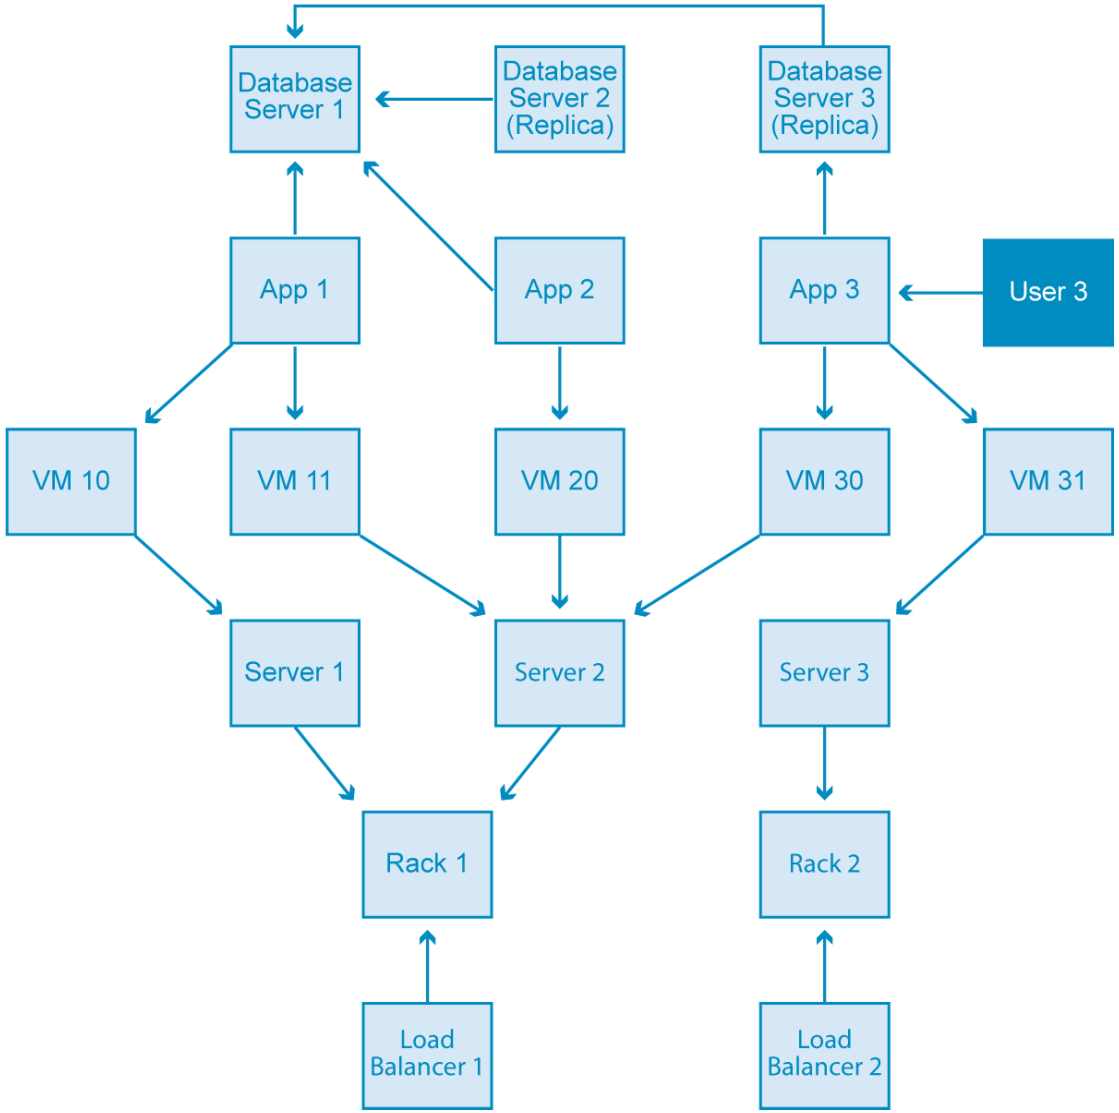

FLORIDA POLYTECHNIC UNIVERSITY

Relational Model

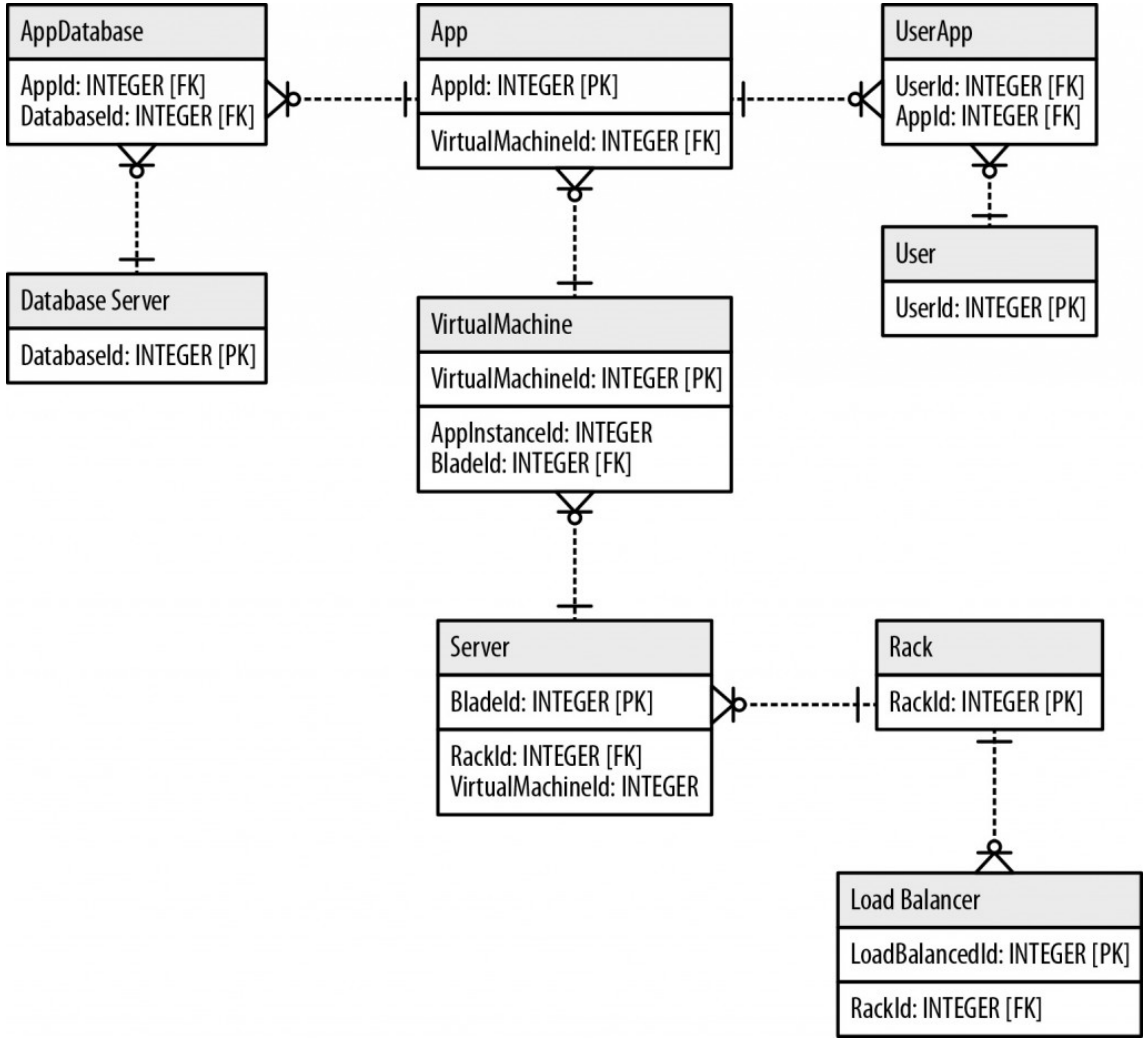

FLORIDA POLYTECHNIC UNIVERSITY

Graph Model

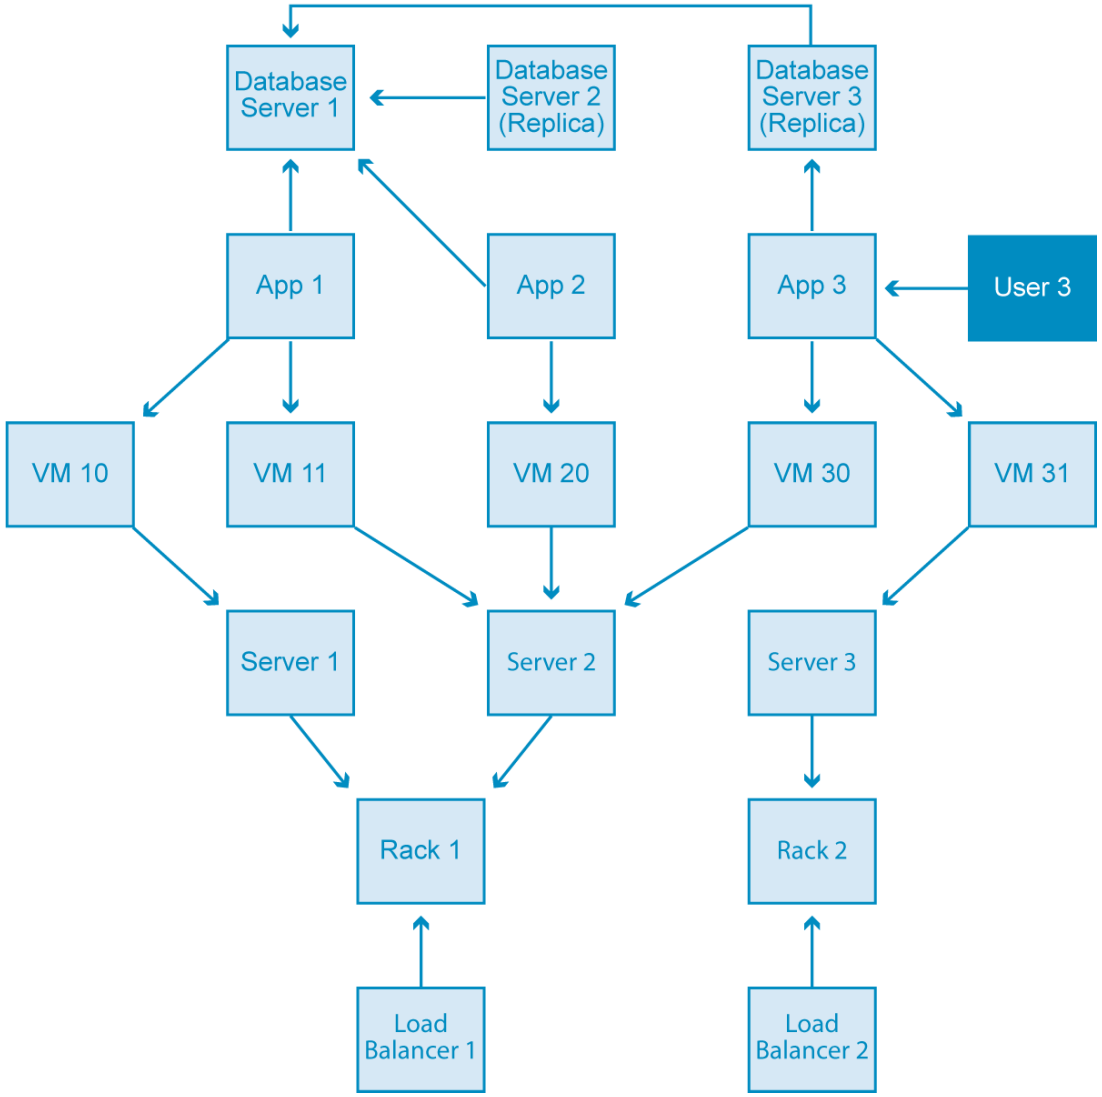

FLORIDA POLYTECHNIC UNIVERSITY

Enriched Model

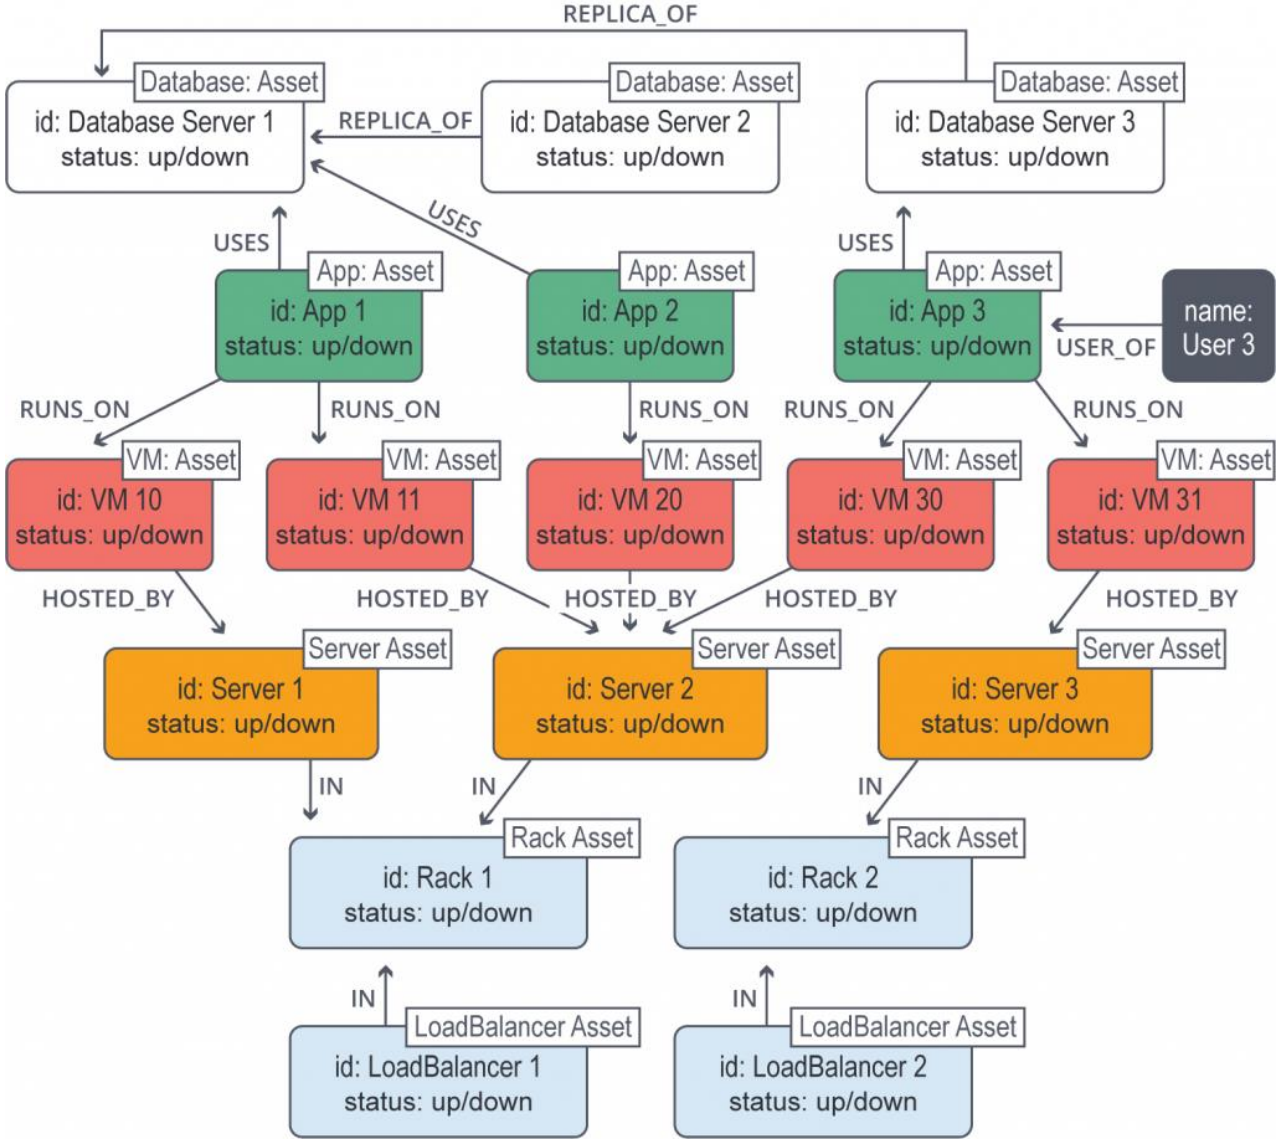

FLORIDA POLYTECHNIC UNIVERSITY

Property Graph Model

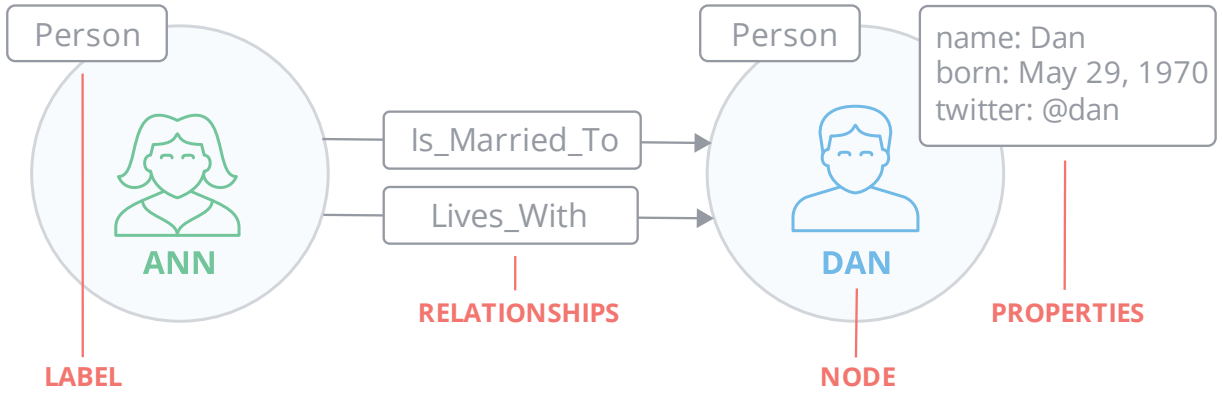

Performance

Depth	RDBMS execution time(s)	Neo4j execution time(s)	Records returned
2	0.016	0.01	~2500
3	30.267	0.168	~110,000
4	1543.505	1.359	~600,000
5	Unfinished	2.132	~800,000

Intelligent Commerce for eBay App on Google Assistant

- Knowledge graph (Neo4j)
- Natural language understanding
- Artificial intelligence
- Store, remember and learn from past interactions with shoppers.
- eBay chose Neo4j as the native graph database that holds the probabilistic models that aid understanding in the conversational shopping scenario.
- The Neo4j graph contains both the product catalog and the attributes of shopper interactions while seeking products.

Examples

- https://neo4j.com/graphgist/degrees-offered-by-the-university-of-oviedo
- https://neo4j.com/graphgist/fitness-and-nutritional-recommendations
- https://neo4j.com/graphgist/credit-card-fraud-detection